

# Compositional Reasoning for WMMs

## COV889 Course Presentation

Ramneet Singh  
Mrunmayi Bhalerao

IIT Delhi

May 2023

- 1 Introduction
- 2 Abstract Language
- 3 Basic Proof System
- 4 Multicopy Atomic Memory Models
- 5 Non-Multicopy Atomic Memory Models

## Section 1

### Introduction

# Rely/Guarantee Reasoning

- Verification of concurrent programs with shared resources is challenging due to **combinatorial explosion**
- **Abstraction to the rescue!**
- Everything the environment can do:  $\mathcal{R}$
- Everything you can do:  $\mathcal{G}$

$$\mathcal{R}, \mathcal{G} \vdash P\{c\}Q$$

- *Compositional!*

# Extension to Weak Memory Models

- Judgements using earlier techniques are valid under **sequentially consistent semantics**
  - Can be directly used for data-race free code executing on weak memory models
  - But, lots of code has data races! `seqlock`, `ConcurrentLinkedQueue` in `java.util.concurrent` ...
- How do we extend them to weak memory models?
- What If: We could find a condition under which sequentially consistent rely-guarantee reasoning can be *soundly* preserved

$$(\vdash P\{c\}Q) \wedge ?? \implies \vdash P\{c_{WM}\}Q$$

- Benefits:
  - Reuse existing verification techniques
  - Deal with the complexity of weak memory separately as a *side-condition*

# What are WMMs anyway?

- Relaxing the memory consistency guarantees provided by hardware enables optimisations
  - Store forwarding (will see later)
  - Write buffers
- (Part 1) **Multicopy Atomic**: One thread's stores become observable to all other threads at the same time.
  - x86-TSO, ARMv8, RISC-V
- (Part 2) **Non-Multicopy Atomic**: Each component has its own *view* of the global memory.
  - Older ARM versions, POWER, C11
- Challenge: Two types of interference now – **Inter-Thread** + **Intra-Thread** (due to reordering)
- How will we deal with this? ...

# Teaser

- We want a compositional approach through thread-local reasoning.
- Exploit the reordering semantics of Colvin and Smith: **multicopy atomic memory models can be captured in terms of instruction reordering.**
  - Combinatorial explosion? ( $n$  reorderable instructions in a thread  $\implies n!$  behaviours)
  - Introduce *reordering interference freedom* between  $(\frac{n(n-1)}{2})$  pairs of instructions (Stay tuned...)
- In non-multicopy atomic WMMs, there is **no global shared state(!)**
  - Judgement for each thread is applicable to *its view* (depends on *propagation* of writes by hardware)
  - How do we know it holds in other threads' views?
  - Represent the semantics using reordering **between** different threads
  - No longer compositional? Hardest part of the talk – *global reordering interference freedom*: use the rely abstraction to represent reorderings between threads

## Section 2

### Abstract Language



# Syntax

- Individual (atomic) instructions  $\alpha$
- Commands (or programs)

$$c := \epsilon \mid \alpha \mid c_1; c_2 \mid c_1 \sqcap c_2 \mid c^* \mid c_1 \parallel c_2$$

- Iteration, choice are non-deterministic
- Empty program  $\epsilon$  represents termination

# Semantics: Commands

- Each atomic instruction  $\alpha$  has a relation  $\text{beh}(\alpha)$  (over pre- and post-states) specifying its behaviour
- Program execution is defined by a small-step semantics over commands
- Iteration, non-deterministic choice are dealt with at a higher level (see next slide)

$$\frac{}{\alpha \mapsto_{\alpha} \epsilon} \qquad \frac{c_1 \mapsto_{\alpha} c'_1}{c_1; c_2 \mapsto_{\alpha} c'_1; c_2}$$
$$\frac{c_1 \mapsto_{\alpha} c'_1}{c_1 \parallel c_2 \mapsto_{\alpha} c'_1 \parallel c_2} \qquad \frac{c_2 \mapsto_{\alpha} c'_2}{c_1 \parallel c_2 \mapsto_{\alpha} c_1 \parallel c'_2}$$

# Semantics: Configurations

- *Configuration*  $(c, \sigma)$  of a program
  - Command  $c$  to be executed
  - State  $\sigma$  (map from variables to values)
- *Action Step*: Performed by component, changes state

$$(c, \sigma) \xrightarrow{as} (c', \sigma') \iff \exists \alpha. c \mapsto_{\alpha} c' \wedge (\sigma, \sigma') \in \text{beh}(\alpha)$$

- *Silent Step*: Performed by component, doesn't change state

$$(c_1 \sqcap c_2, \sigma) \rightsquigarrow (c_1, \sigma) \quad (c_1 \sqcap c_2, \sigma) \rightsquigarrow (c_2, \sigma)$$

$$(c^*, \sigma) \rightsquigarrow (\epsilon, \sigma) \quad (c^*, \sigma) \rightsquigarrow (c; c^*, \sigma)$$

- *Program Step*: Action Step or Silent Step
- *Environment Step*: Performed by environment, changes state.  
 $(c, \sigma) \xrightarrow{es} (c, \sigma')$

## Section 3

### Basic Proof System

# Definitions

- Associate a verification condition  $vc(\alpha)$  with each instruction  $\alpha$ :  
Provides finer-grained control (just set to  $\top$  if not needed)

- Hoare triple

$$P\{ \alpha \}Q \stackrel{\text{def}}{=} P \subseteq vc(\alpha) \cap \{ \sigma \mid \forall \sigma', (\sigma, \sigma') \in \text{beh}(\alpha) \implies \sigma' \in Q \}$$

- A rely-guarantee pair  $(\mathcal{R}, \mathcal{G})$  is well-formed if

- $\mathcal{R}$  is reflexive and transitive
- $\mathcal{G}$  is reflexive

- Stability of predicate  $P$  under rely condition  $\mathcal{R}$

$$\text{stable}_{\mathcal{R}}(P) \stackrel{\text{def}}{=} P \subseteq \{ \sigma \in P \mid \forall \sigma', (\sigma, \sigma') \in \mathcal{R} \implies \sigma' \in P \}$$

- Instruction  $\alpha$  satisfies guarantee condition  $\mathcal{G}$

$$\text{sat}(\alpha, \mathcal{G}) \stackrel{\text{def}}{=} \{ \sigma \mid \forall \sigma', (\sigma, \sigma') \in \text{beh}(\alpha) \implies (\sigma, \sigma') \in \mathcal{G} \}$$

- Now introduce rely/guarantee judgements at three levels

# Instruction Level ( $\vdash_a$ )

$$\mathcal{R}, \mathcal{G} \vdash_a P\{ \alpha \} Q \stackrel{\text{def}}{=} \text{stable}_{\mathcal{R}}(P) \wedge \text{stable}_{\mathcal{R}}(Q) \wedge \text{vc}(\alpha) \subseteq \text{sat}(\alpha, \mathcal{G}) \wedge P\{ \alpha \} Q$$

- Interplay between environmental interference and pre-, post-conditions handled through stability

# Component Level ( $\vdash_c$ )

$$\text{Atom} \frac{\mathcal{R}, \mathcal{G} \vdash_a P\{\alpha\}Q}{\mathcal{R}, \mathcal{G} \vdash_c P\{\alpha\}Q}$$

$$\text{Seq} \frac{\mathcal{R}, \mathcal{G} \vdash_c P\{c_1\}M \quad \mathcal{R}, \mathcal{G} \vdash_c M\{c_2\}Q}{\mathcal{R}, \mathcal{G} \vdash_c P\{c_1; c_2\}Q}$$

$$\text{Choice} \frac{\mathcal{R}, \mathcal{G} \vdash_c P\{c_1\}Q \quad \mathcal{R}, \mathcal{G} \vdash_c P\{c_2\}Q}{\mathcal{R}, \mathcal{G} \vdash_c P\{c_1 \sqcap c_2\}Q}$$

$$\text{Iteration} \frac{\mathcal{R}, \mathcal{G} \vdash_c P\{c\}P \quad \text{stable}_{\mathcal{R}}(P)}{\mathcal{R}, \mathcal{G} \vdash_c P\{c^*\}P}$$

$$\text{Conseq} \frac{\mathcal{R}, \mathcal{G} \vdash_c P\{c\}Q \quad P' \subseteq P \quad \mathcal{R}' \subseteq \mathcal{R} \quad Q \subseteq Q' \quad \mathcal{G} \subseteq \mathcal{G}'}{\mathcal{R}', \mathcal{G}' \vdash_c P'\{c\}Q'}$$

# Global Level ( $\vdash$ )

- Global satisfiability needs component satisfiability + **interference check**

$$\text{Comp} \frac{\mathcal{R}, \mathcal{G} \vdash_c P\{c\}Q \quad \text{rif}(\mathcal{R}, \mathcal{G}, c)}{\mathcal{R}, \mathcal{G} \vdash P\{c\}Q}$$

- Usual parallel rule

$$\text{Par} \frac{\mathcal{R}_1, \mathcal{G}_1 \vdash_c P_1\{c_1\}Q_1 \quad \mathcal{R}_2, \mathcal{G}_2 \vdash_c P_2\{c_2\}Q_2 \quad \mathcal{G}_2 \subseteq \mathcal{R}_1 \quad \mathcal{G}_1 \subseteq \mathcal{R}_2}{\mathcal{R}_1 \cap \mathcal{R}_2, \mathcal{G}_1 \cup \mathcal{G}_2 \vdash P_1 \wedge P_2\{c_1 \parallel c_2\}Q_1 \wedge Q_2}$$



## Section 4

# Multicopy Atomic Memory Models

# Reordering Semantics: Basics

- Multicopy atomic memory models can be characterised using a *reordering* relation  $\leftarrow$  over pairs of instructions in a component
- $\leftarrow$  is syntactically derivable based on the specific memory model.  
E.g., in ARMv8
  - Two instructions which don't access (read or write) a common variable can be reordered
  - Various types of memory barriers prevent reordering
- *Forwarding* is another complication
  - $\beta = x := 3; \alpha = y := x$ . Can forward the value 3 to  $y$ , losing dependence between  $\alpha, \beta$ .
  - $x := 3 ; y := x \implies y := 3 ; x := 3$
  - Denote  $\alpha$  with the value written in an earlier instruction forwarded to it as  $\alpha_{<\beta>}$ .
- Forwarding may continue arbitrarily and can span multiple instructions

# Reordering Semantics: Formal

- $\alpha_{<c>}$ : cumulative forwarding effects of the instructions in command  $c$  on  $\alpha$
- Ternary relation  $\gamma < c < \alpha$ : Reordering of instruction  $\alpha$  prior to command  $c$ , with cumulative forwarding effects producing  $\gamma$ .
- Definition by induction

$$\alpha_{<\beta>} < \beta < \alpha \stackrel{\text{def}}{=} \beta \leftrightarrow \alpha_{<\beta>}$$

$$\alpha_{<c_1; c_2>} < c_1; c_2 < \alpha \stackrel{\text{def}}{=} \alpha_{<c_1; c_2>} < c_1 < \alpha_{<c_2>} \wedge \alpha_{<c_2>} < c_2 < \alpha$$

- Example:  $\alpha = (y := x), \beta = (x := 3), \gamma = (z := 5)$ .  
 $\alpha_{<\beta>} = (y := 3), \alpha_{<\gamma; \beta>} = (y := 3)$ .

$y := 3 < x := 3 < y := x$  and  $y := 3 < z := 5 ; x := 3 < y := x$

- Can execute an instruction which occurs later in the program if reordering and forwarding can bring it (in its new form  $\gamma$ ) to the beginning

$$\text{Reorder} \frac{c_2 \mapsto_{\alpha} \quad \gamma < c < \alpha}{c_1; c_2 \mapsto_{\gamma} c_1; c'_2}$$

# Reordering Interference Freedom

- Insight: Any valid reordering will **preserve thread-local semantics**, thus may only invalidate reasoning when **observed by the environment**.
  - Abstraction to the rescue again! Observed by environment  $\implies \mathcal{G}$  violated, or  $\mathcal{R}$  not strong enough
- Three Levels: Instructions, Commands, Program

# RIF: Instructions

- Two instructions are *reordering interference free*: Reasoning over them in their original order is sufficient to include reordered behaviour.

$$\begin{aligned} \text{rif}_a(\mathcal{R}, \mathcal{G}, \beta, \alpha) &\stackrel{\text{def}}{=} \forall P, Q, M. \mathcal{R}, \mathcal{G} \vdash_a P\{\beta\}M \wedge \mathcal{R}, \mathcal{G} \vdash_a M\{\alpha\}Q \\ &\implies \exists M'. \mathcal{R}, \mathcal{G} \vdash_a P\{\alpha_{<\beta>}\}M' \wedge \mathcal{R}, \mathcal{G} \vdash_a M'\{\beta\}Q \end{aligned}$$

# RIF: Commands

- Command  $c$  is *reordering interference free* from  $\alpha$  under  $\mathcal{R}, \mathcal{G}$  if the reordering of  $\alpha$  over each instruction of  $c$  is reordering interference free, *including those variants produced by forwarding*.

$$\text{rif}_c(\mathcal{R}, \mathcal{G}, \beta, \alpha) \stackrel{\text{def}}{=} \text{rif}_a(\mathcal{R}, \mathcal{G}, \beta, \alpha)$$

$$\text{rif}_c(\mathcal{R}, \mathcal{G}, c_1; c_2, \alpha) \stackrel{\text{def}}{=} \text{rif}_c(\mathcal{R}, \mathcal{G}, c_1, \alpha_{<c_2>}) \wedge \text{rif}_c(\mathcal{R}, \mathcal{G}, c_2, \alpha)$$

# RIF: Programs

- Program  $c$  is *reordering interference free* if and only if **all possible reorderings** of its instructions over the respective prefixes are reordering interference free.

$$\text{rif}(\mathcal{R}, \mathcal{G}, c) \stackrel{\text{def}}{=} \forall \alpha, r, c'. c \mapsto_{\alpha <_r} c' \implies \text{rif}_c(\mathcal{R}, \mathcal{G}, r, \alpha) \wedge \text{rif}(\mathcal{R}, \mathcal{G}, c')$$

- Observe: Checking  $\text{rif}(\mathcal{R}, \mathcal{G}, c)$  amounts to
  - Checking  $\text{rif}_a(\mathcal{R}, \mathcal{G}, \beta, \alpha)$  for all pairs of instructions  $\beta, \alpha$  that can reorder in  $c$
  - Including those pairs for which  $\alpha$  is a new instruction generated through forwarding

# Gameplan

- ① Compute all pairs of reorderable instructions  $(\beta, \alpha)$ .
- ② Demonstrate reordering interference freedom for as many of these pairs as possible (using  $\text{rif}_a(\mathcal{R}, \mathcal{G}, \beta, \alpha)$ ).
- ③ If  $\text{rif}_a$  cannot be shown for some pairs
  - introduce memory barriers to prevent their reordering or
  - modify the verification problem such that their reordering can be considered benign
- ④ Verify the component in isolation, using standard rely/guarantee reasoning with an assumed sequentially consistent memory model.

For a thread with  $n$  reorderable instructions,

$n!$  Possible Behaviours  $\longrightarrow n(n-1)/2$   $\text{rif}_a$  checks

Thanks for staying tuned : )



## Section 5

# Non-Multicopy Atomic Memory Models

# Non-Multicopy Atomic Memory Models

- There is **no shared state** that all components agree on throughout execution, invalidating a core assumption of standard rely/guarantee reasoning.

# Write History Semantics: Representation

- Each component is associated with a unique identifier.
- Shared memory state is represented as a list of variable writes  $\langle w_1, w_2, w_3, \dots \rangle$ , with metadata to indicate which components have performed and observed particular writes.
- The order of events in this write history provides an *overall order* to the system's events, with those later in the list being the most recent.
- Each  $w_i = (x \mapsto v)_{rds}^{wr}$  where
  - $x$  is a variable
  - $v$  is a value
  - $\text{writer}((x \mapsto v)_{rds}^{wr}) = wr$  is the writer component's identifier
  - $\text{readers}((x \mapsto v)_{rds}^{wr}) = rds$  is the set of component identifiers that have observed the write
  - $\text{var}((x \mapsto v)_{rds}^{wr}) = x$

# Write History Semantics: Manipulation

- Divide instructions into two types: *global* and *local*. Global instructions  $\alpha$  are:
  - Store  $(x := v)_i$ , Load  $[x = v]_i$ , Memory barrier fence $_i$ , Skip instruction (corresponding to some internal step)
- Behaviour of these instructions is formalised as (for skip it's just *id*):

$$\text{beh}((x := v)_i) = \{ (h \circ h', h \circ (x \mapsto v)_{\{i\}}^i \circ h') \mid$$

$$\forall w \in h'. \text{writer}(w) \neq i \wedge (\text{var}(w) = x \implies i \notin \text{readers}(w))$$

$$\text{beh}([x = v]_i) = \{ (h \circ (x \mapsto v)_r^j \circ h', h \circ (x \mapsto v)_r^j \circ h') \mid$$

$$\forall w \in h'. (\text{var}(w) = x \implies i \notin \text{readers}(w))$$

$$\text{beh}(\text{fence}_i) = \{ (h, h) \mid \forall w \in h. (i \in \text{readers}(w) \implies \forall y. y \in \text{readers}(w))$$

- Propagations of writes are modelled as environment effects and can take place at any point during the execution.

$$\text{prp} = \{ (h \circ (x \mapsto v)_r^j \circ h', h \circ (x \mapsto v)_{r \cup \{i\}}^j \circ h') \mid$$

$$i \notin r \wedge \forall w \in h. (\text{var}(w) = x \implies i \in \text{readers}(w)) \}$$

# More Notation

- New constructor in the language:  $\text{comp}(i, m, c)$  indicating a component with
  - identifier  $i$
  - local state  $m$
  - command  $c$
- Assume a local behaviour relation  $\text{lbeh}$  such that  $(m, \alpha', m') \in \text{lbeh}(\alpha)$  if executing  $\alpha$ 
  - changes the local state from  $m$  to  $m'$
  - corresponds to the global instruction  $\alpha'$

$$\text{comp}(i, m, c) \mapsto_{\alpha'_i} \text{comp}(i, m', c') \iff c \mapsto_{\alpha} c' \wedge (m, \alpha', m') \in \text{lbeh}(\alpha)$$

- Go from local semantics/reasoning to global semantics/reasoning using  $\text{comp}$  and  $\text{lbeh}$ .
  - Constraint: systems are constructed as the parallel composition of a series of  $\text{comp}$  commands.
  - Trivial support for local state (e.g., hardware registers).

# Meaning of Judgement

- If there is no global state, what does  $\mathcal{R}, \mathcal{G} \vdash P\{c\}Q$  (for a component  $i$  with command  $c$ ) mean?
- For a set of components  $I$ , write history  $h$ , for all variables  $x$ ,  $\text{view}_I(h, x) = v$  iff

$$h = h' \circ (x \mapsto v)_r^w \circ h'' \wedge I \subseteq r \wedge \forall w_i \in h''. \text{var}(w_i) = x \implies I \not\subseteq \text{readers}(w_i)$$

- For all executions of  $c$ 
  - If
    - the execution operates on a write history  $h$  such that  $\text{view}_i(h) \in P$
    - all propagations to  $i$  modify  $\text{view}_i$  in accordance with  $\mathcal{R}$
  - Then  $i$  will
    - modify  $\text{view}_i$  in accordance with  $\mathcal{G}$
    - given termination, end with a write history  $h$  such that  $\text{view}_i(h) \in Q$
- This state mapping allows for rely/guarantee judgements **over individual components** to be trivially lifted from a standard memory model to their respective views of a write history.

# Parallel Composition: Taming the Beast

- Parallel composition is complicated: **Need to relate differing components views.**
- *If* the execution of an instruction  $\alpha$  by some component  $i$  satisfies its guarantee specification  $\mathcal{G}_i$  in state  $h$ ,

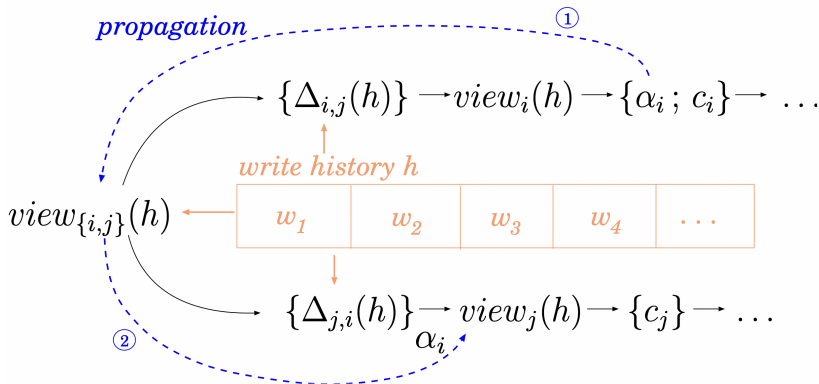
$$\text{view}_i(h) \in \text{sat}(\alpha, \mathcal{G}_i)$$

- *Then* the effects of propagating  $\alpha$ 's writes to some other component  $j$  will satisfy its rely specification  $\mathcal{R}_j$  in its view,

$$\text{view}_j(h) \in \text{sat}(\alpha, \mathcal{R}_j)$$

- Insight: It is possible to relate the views of two components by only considering the **difference in their observed writes**, i.e., the writes one component has observed but the other has not.

# Travelling Between Components



- Aim to demonstrate rely/guarantee compatibility when propagating an instruction  $\alpha$  from component  $i$  to component  $j$ 
  - Given: component  $i$  executes  $\alpha$  such that  $view_i(h) \in \text{sat}(\alpha, \mathcal{G}_i)$ .
  - Step 1: Show that  $\alpha$  can be executed in the shared view, i.e.,  $view_{\{i,j\}}(h) \in \text{sat}(\alpha, \mathcal{G}_i)$ .
  - Step 2: Show that  $\alpha$  can be executed in component  $j$ 's view, i.e.,  $view_j(h) \in \text{sat}(\alpha, \mathcal{G}_i)$ .



# Travelling to Shared View Through Non $i, j$ Writes

- Prove step 1 by induction on length of  $\Delta_{i,j}(h)$ .
  - Base case trivial ( $\text{view}_{\{i,j\}}(h) = \text{view}_i(h)$ ).
  - Induction step:
    - The write cannot be from  $j$  since  $j$  hasn't observed it.
    - If the write is from  $i$ , then the reordering of  $\alpha$  before it has been covered in multicopy reordering interference freedom.
    - If the write is from some  $k \neq i, j$ , then we do what follows.
- Have some relation  $\mathcal{E}$  intended to capture the possible writes  $i$  may have observed ahead of  $j$

$$\text{rif}_{nmca}(\mathcal{E}, \alpha, \mathcal{G}_i) = wp(\mathcal{E}, \text{sat}(\alpha, \mathcal{G}_i)) \subseteq \text{sat}(\alpha, \mathcal{G}_i)$$

- Proving this for  $\mathcal{E} = \mathcal{R}_i \cap \mathcal{R}_j \cap id_\alpha$  is sufficient.

# Endgame

- Define a compatibility relation by universally quantifying over all writes

$$\text{compat}(\mathcal{G}_i, \mathcal{R}_i, \mathcal{R}_j) \stackrel{\text{def}}{=} \forall x, v.$$

$$wp(\mathcal{R}_i \cap \mathcal{R}_j \cap id_x, \text{sat}(x := v, \mathcal{G}_i)) \subseteq \text{sat}(x := v, \mathcal{R}_j)$$

- Modify the rules for parallel composition (note that we need separate relies and guarantees for each component because demonstrating compat requires pairwise checking)

$$\text{Comp}' \frac{\mathcal{R}, \mathcal{G} \vdash_c P\{c\}Q \quad \text{rif}(\mathcal{R}, \mathcal{G}, c)}{[i \mapsto \mathcal{R}], [i \mapsto \mathcal{G}] \vdash P\{\text{comp}(i, m, c)\}Q}$$

$$\text{Par}' \frac{\begin{array}{l} \mathcal{R}_1, \mathcal{G}_1 \vdash P_1\{c_1\}Q_1 \quad \mathcal{R}_2, \mathcal{G}_2 \vdash P_2\{c_2\}Q_2 \quad \text{disjoint}(\mathcal{R}_1, \mathcal{R}_2) \\ \forall i \in \text{dom}(\mathcal{R}_1). \forall j \in \text{dom}(\mathcal{R}_2). \text{compat}(\mathcal{G}_1(i), \mathcal{R}_1(i), \mathcal{R}_2(j)) \\ \forall i \in \text{dom}(\mathcal{R}_2). \forall j \in \text{dom}(\mathcal{R}_1). \text{compat}(\mathcal{G}_2(i), \mathcal{R}_2(i), \mathcal{R}_1(j)) \end{array}}{\mathcal{R}_1 \uplus \mathcal{R}_2, \mathcal{G}_1 \uplus \mathcal{G}_2 \vdash P_1 \wedge P_2\{c_1 \parallel c_2\}Q_1 \wedge Q_2}$$