

Compositional Reasoning for WMMs

COV889 Course Presentation

Ramneet Singh
Mrunmayi Bhalerao

IIT Delhi

May 2023

- 1 Introduction
- 2 Abstract Language
- 3 Basic Proof System
- 4 Multicopy Atomic Memory Models

Section 1

Introduction

Rely/Guarantee Reasoning

- Verification of concurrent programs with shared resources is challenging due to **combinatorial explosion**
- **Abstraction to the rescue!**
- Everything the environment can do: \mathcal{R}
- Everything you can do: \mathcal{G}

$$\mathcal{R}, \mathcal{G} \vdash P\{c\}Q$$

- *Compositional!*

Extension to Weak Memory Models

- Judgements using earlier techniques are valid under **sequentially consistent semantics**
 - Can be directly used for data-race free code executing on weak memory models
 - But, lots of code has data races! `seqlock`,
`java.util.concurrent.ConcurrentLinkedQueue` ...
- How do we extend them to weak memory models?
- What If: We could find a condition under which sequentially consistent rely-guarantee reasoning can be *soundly* preserved

$$(\vdash P\{c\}Q) \wedge ?? \implies \vdash P\{c_{WM}\}Q$$

- Benefits:
 - Reuse existing verification techniques
 - Deal with the complexity of weak memory separately as a *side-condition*

What are WMMs anyway?

- Relaxing the memory consistency guarantees provided by hardware enables optimisations
 - Store forwarding (will see later)
 - Write buffers
- (Part 1) **Multicopy Atomic**: One thread's stores become observable to all other threads at the same time.
 - x86-TSO, ARMv8, RISC-V
- (Part 2) **Non-Multicopy Atomic**: Each component has its own *view* of the global memory.
 - Older ARM versions, POWER, C11
- Challenge: Two types of interference now – **Inter-Thread** + **Intra-Thread** (due to reordering)
- How will we deal with this? ...

Teaser

- We want a compositional approach through thread-local reasoning.
- Exploit the reordering semantics of Colvin and Smith: **multicopy atomic memory models can be captured in terms of instruction reordering.**
 - Combinatorial explosion? (n reorderable instructions in a thread $\implies n!$ behaviours)
 - Introduce *reordering interference freedom* between $(\frac{n(n-1)}{2})$ pairs of instructions (Stay tuned. . .)
- In non-multicopy atomic WMMs, there is **no global shared state**(!!)
 - Judgement for each thread is applicable to *its view* (depends on *propagation* of writes by hardware)
 - How do we know it holds in other threads' views?
 - Represent the semantics using reordering **between** different threads
 - No longer compositional? Hardest part of the talk – *global reordering interference freedom*: use the rely abstraction to represent reorderings between threads

Section 2

Abstract Language

Syntax

- Individual (atomic) instructions α
- Commands (or programs)

$$c := \epsilon \mid \alpha \mid c_1; c_2 \mid c_1 \sqcap c_2 \mid c^* \mid c_1 \parallel c_2$$

- Iteration, choice are non-deterministic
- Empty program ϵ represents termination

Semantics: Commands

- Each atomic instruction α has a relation $\text{beh}(\alpha)$ (over pre- and post-states) specifying its behaviour
- Program execution is defined by a small-step semantics over commands
- Iteration, non-deterministic choice are dealt with at a higher level (see next slide)

$$\begin{array}{c}
 \frac{}{\alpha \mapsto_{\alpha} \epsilon}
 \end{array}
 \qquad
 \frac{c_1 \mapsto_{\alpha} c'_1}{c_1; c_2 \mapsto_{\alpha} c'_1; c_2}$$

$$\frac{c_1 \mapsto_{\alpha} c'_1}{c_1 \parallel c_2 \mapsto_{\alpha} c'_1 \parallel c_2}
 \qquad
 \frac{c_2 \mapsto_{\alpha} c'_2}{c_1 \parallel c_2 \mapsto_{\alpha} c_1 \parallel c'_2}$$

Semantics: Configurations

- *Configuration* (c, σ) of a program
 - Command c to be executed
 - State σ (map from variables to values)
- *Action Step*: Performed by component, changes state

$$(c, \sigma) \xrightarrow{as} (c', \sigma') \iff \exists \alpha. c \mapsto_{\alpha} c' \wedge (\sigma, \sigma') \in \text{beh}(\alpha)$$

- *Silent Step*: Performed by component, doesn't change state

$$(c_1 \sqcap c_2, \sigma) \rightsquigarrow (c_1, \sigma) \quad (c_1 \sqcap c_2, \sigma) \rightsquigarrow (c_2, \sigma)$$

$$(c^*, \sigma) \rightsquigarrow (\epsilon, \sigma) \quad (c^*, \sigma) \rightsquigarrow (c; c^*, \sigma)$$

- *Program Step*: Action Step or Silent Step
- *Environment Step*: Performed by environment, changes state.
 $(c, \sigma) \xrightarrow{es} (c, \sigma')$

Section 3

Basic Proof System

Definitions

- Associate a verification condition $vc(\alpha)$ with each instruction α :
Provides finer-grained control (just set to \top if not needed)

- Hoare triple

$$P\{ \alpha \}Q \stackrel{\text{def}}{=} P \subseteq vc(\alpha) \cap \{ \sigma \mid \forall \sigma', (\sigma, \sigma') \in \text{beh}(\alpha) \implies \sigma' \in Q \}$$

- A rely-guarantee pair $(\mathcal{R}, \mathcal{G})$ is well-formed if

- \mathcal{R} is reflexive and transitive
- \mathcal{G} is reflexive

- Stability of predicate P under rely condition \mathcal{R}

$$\text{stable}_{\mathcal{R}}(P) \stackrel{\text{def}}{=} P \subseteq \{ \sigma \in P \mid \forall \sigma', (\sigma, \sigma') \in \mathcal{R} \implies \sigma' \in P \}$$

- Instruction α satisfies guarantee condition \mathcal{G}

$$\text{sat}(\alpha, \mathcal{G}) \stackrel{\text{def}}{=} \{ \sigma \mid \forall \sigma', (\sigma, \sigma') \in \text{beh}(\alpha) \implies (\sigma, \sigma') \in \mathcal{G} \}$$

- Now introduce rely/guarantee judgements at three levels

Instruction Level (\vdash_a)

$$\mathcal{R}, \mathcal{G} \vdash_a P\{c\}Q \stackrel{\text{def}}{=} \text{stable}_{\mathcal{R}}(P) \wedge \text{stable}_{\mathcal{R}}(Q) \wedge \text{vc}(\alpha) \subseteq \text{sat}(\alpha, \mathcal{G}) \wedge P\{c\}Q$$

- Interplay between environmental interference and pre-, post-conditions handled through stability

Component Level (\vdash_c)

$$\text{Atom} \frac{\mathcal{R}, \mathcal{G} \vdash_a P\{\alpha\}Q}{\mathcal{R}, \mathcal{G} \vdash_c P\{\alpha\}Q}$$

$$\text{Seq} \frac{\mathcal{R}, \mathcal{G} \vdash_c P\{c_1\}M \quad \mathcal{R}, \mathcal{G} \vdash_c M\{c_2\}Q}{\mathcal{R}, \mathcal{G} \vdash_c P\{c_1; c_2\}Q}$$

$$\text{Choice} \frac{\mathcal{R}, \mathcal{G} \vdash_c P\{c_1\}Q \quad \mathcal{R}, \mathcal{G} \vdash_c P\{c_2\}Q}{\mathcal{R}, \mathcal{G} \vdash_c P\{c_1 \sqcap c_2\}Q}$$

$$\text{Iteration} \frac{\mathcal{R}, \mathcal{G} \vdash_c P\{c\}P \quad \text{stable}_{\mathcal{R}}(P)}{\mathcal{R}, \mathcal{G} \vdash_c P\{c^*\}Q}$$

$$\text{Conseq} \frac{\mathcal{R}, \mathcal{G} \vdash_c P\{c\}Q \quad P' \subseteq P \quad \mathcal{R}' \subseteq \mathcal{R} \quad Q \subseteq Q' \quad \mathcal{G} \subseteq \mathcal{G}'}{\mathcal{R}', \mathcal{G}' \vdash_c P'\{c\}Q'}$$

Global Level (\vdash)

- Global satisfiability needs component satisfiability + **interference check**

$$\text{Comp} \frac{\mathcal{R}, \mathcal{G} \vdash_c P\{c\}Q \quad \text{rif}(\mathcal{R}, \mathcal{G}, c)}{\mathcal{R}, \mathcal{G} \vdash P\{c\}Q}$$

- Usual parallel rule

$$\text{Par} \frac{\mathcal{R}, \mathcal{G} \vdash_c P\{c\}Q \quad \text{rif}(\mathcal{R}, \mathcal{G}, c)}{\mathcal{R}, \mathcal{G} \vdash P\{c\}Q}$$

Section 4

Multicopy Atomic Memory Models

Reordering Semantics: Basics

- Multicopy atomic memory models can be characterised using a *reordering* relation \leftarrow over pairs of instructions in a component
- \leftarrow is syntactically derivable based on the specific memory model.
E.g., in ARMv8
 - Two instructions which don't access (read or write) a common variable can be reordered
 - Various types of memory barriers prevent reordering
- *Forwarding* is another complication
 - $\beta = x := 3; \alpha = y := x$. Can forward the value 3 to y , losing dependence between α, β .
 - $x := 3 ; y := x \implies y := 3 ; x := 3$
 - Denote α with the value written in an earlier instruction forwarded to it as $\alpha_{<\beta>}$.
- Forwarding may continue arbitrarily and can span multiple instructions

Reordering Semantics: Formal

- $\alpha_{<c>}$: cumulative forwarding effects of the instructions in command c on α
- Ternary relation $\gamma < c < \alpha$: Reordering of instruction α prior to command c , with cumulative forwarding effects producing γ .
- Definition by induction

$$\alpha_{<\beta>} < \beta < \alpha \stackrel{\text{def}}{=} \beta \leftrightarrow \alpha_{<\beta>}$$

$$\alpha_{<c_1; c_2>} < c_1; c_2 < \alpha \stackrel{\text{def}}{=} \alpha_{<c_1; c_2>} < c_1 < \alpha_{<c_2>} \wedge \alpha_{<c_2>} < c_2 < \alpha$$

- Example: $\alpha = (y := x), \beta = (x := 3), \gamma = (z := 5)$.

$$\alpha_{<\beta>} = (y := 3), \alpha_{<\gamma; \beta>} = (y := 3).$$

$$y := 3 < x := 3 < y := x \quad y := 3 < z := 5 ; x := 3 < y := x$$

- Can execute an instruction which occurs later in the program if reordering and forwarding can bring it (in its new form γ) to the beginning

$$\text{Reorder} \frac{c_2 \mapsto_{\alpha} \quad \gamma < c < \alpha}{c_1; c_2 \mapsto_{\gamma} c_1; c'_2}$$

Reordering Interference Freedom

- Insight: Any valid reordering will **preserve thread-local semantics**, thus may only invalidate reasoning when **observed by the environment**.
 - Abstraction to the rescue again! Observed by environment $\implies \mathcal{G}$ violated, or \mathcal{R} not strong enough
- Three Levels: Instructions, Commands, Program

RIF: Instructions

- Two instructions are *reordering interference free*: Reasoning over them in their original order is sufficient to include reordered behaviour.

$$\begin{aligned} \text{rif}_a(\mathcal{R}, \mathcal{G}, \beta, \alpha) &\stackrel{\text{def}}{=} \forall P, Q, M. \mathcal{R}, \mathcal{G} \vdash_a P\{\beta\}M \wedge \mathcal{R}, \mathcal{G} \vdash_a M\{\alpha\}Q \\ &\implies \exists M'. \mathcal{R}, \mathcal{G} \vdash_a P\{\alpha_{<\beta>}\}M' \wedge \mathcal{R}, \mathcal{G} \vdash_a M'\{\beta\}Q \end{aligned}$$

RIF: Commands

- Command c is *reordering interference free* from α under \mathcal{R}, \mathcal{G} if the reordering of α over each instruction of c is reordering interference free, *including those variants produced by forwarding*.

$$\text{rif}_c(\mathcal{R}, \mathcal{G}, \beta, \alpha) \stackrel{\text{def}}{=} \text{rif}_a(\mathcal{R}, \mathcal{G}, \beta, \alpha)$$

$$\text{rif}_c(\mathcal{R}, \mathcal{G}, c_1; c_2, \alpha) \stackrel{\text{def}}{=} \text{rif}_c(\mathcal{R}, \mathcal{G}, c_1, \alpha_{<c_2>}) \wedge \text{rif}_c(\mathcal{R}, \mathcal{G}, c_2, \alpha)$$

RIF: Programs

- Program c is *reordering interference free* if and only if **all possible reorderings** of its instructions over the respective prefixes are reordering interference free.

$$\text{rif}(\mathcal{R}, \mathcal{G}, c) \stackrel{\text{def}}{=} \forall \alpha, r, c'. c \mapsto_{\alpha_{<r}} c' \implies \text{rif}_c(\mathcal{R}, \mathcal{G}, r, \alpha) \wedge \text{rif}(\mathcal{R}, \mathcal{G}, c')$$

- Observe: Checking $\text{rif}(\mathcal{R}, \mathcal{G}, c)$ amounts to
 - Checking $\text{rif}_a(\mathcal{R}, \mathcal{G}, \beta, \alpha)$ for all pairs of instructions β, α that can reorder in c
 - Including those pairs for which α is a new instruction generated through forwarding

Gameplan

- ① Compute all pairs of reorderable instructions (β, α) .
- ② Demonstrate reordering interference freedom for as many of these pairs as possible (using $\text{rif}_a(R, G, \beta, \alpha)$).
- ③ If rif_a cannot be shown for some pairs
 - introduce memory barriers to prevent their reordering or
 - modify the verification problem such that their reordering can be considered benign
- ④ Verify the component in isolation, using standard rely/guarantee reasoning with an assumed sequentially consistent memory model.

For a thread with n reorderable instructions,

$n!$ Possible Behaviours $\longrightarrow n(n-1)/2$ rif_a checks

Thanks for staying tuned :)